

# Compiler-Aided Optimization

Andrey Kamaev  
Itseez UNN 2016

# What is optimization?

*Finding an **alternative** with the most cost effective or highest achievable **performance** under the given **constraints**, by maximizing desired factors and minimizing undesired ones.*

# Where get the performance?

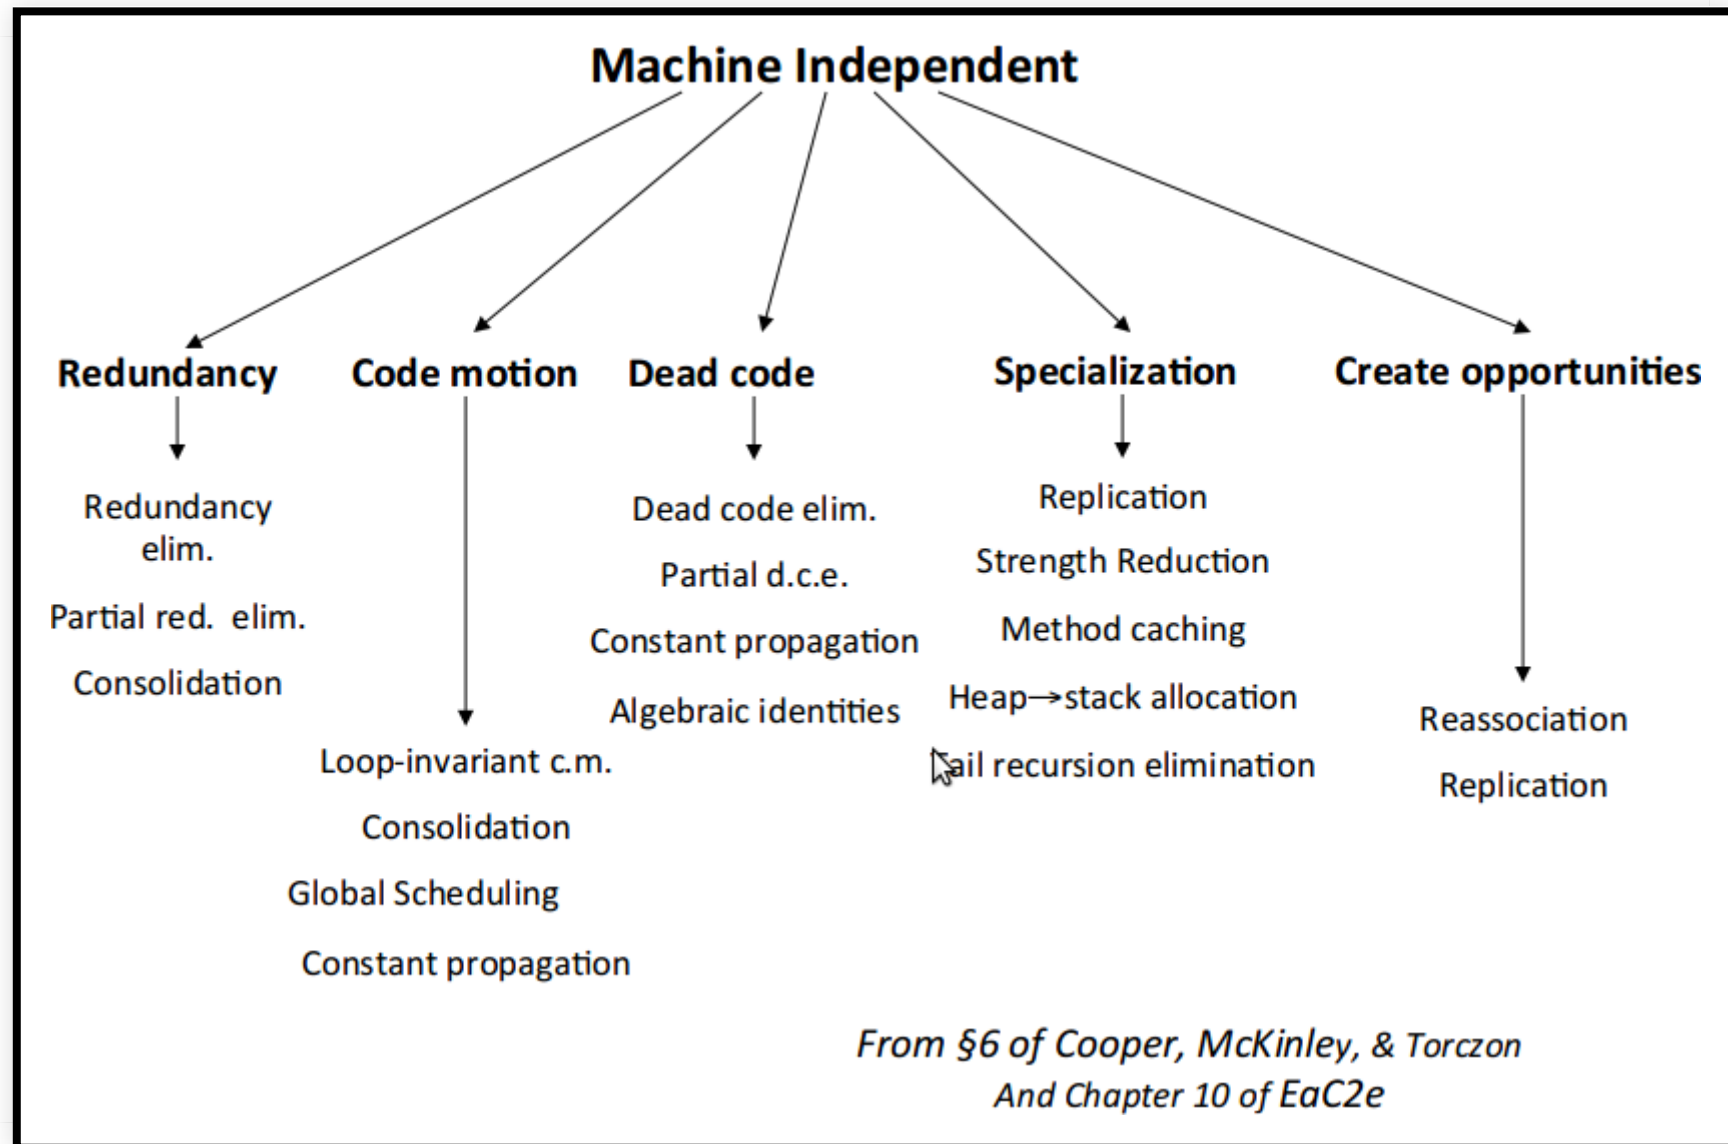
High-level	Programmer
<b>Middle-level</b>	<b>Compiler</b>
Low-level	Hardware

# Understanding Compilers

- Compiler has limited understanding of the program's behavior and the environment in which it will be used
  - Most analysis is performed only within procedures
  - Most analysis is based only on static information
- Compilers emphasize **correctness** rather than **performance**
  - Must not cause any change in program behavior under any possible condition
- On well recognized constructs, compilers will usually do better than the developer
- Often, by simply slightly reorganizing existing code, it is possible to improve both code size and speed

# Classes of compiler optimizations

- High-level optimizations
  - Loop optimizations
  - Interprocedural optimization
- Global optimizations
- Local optimizations
- Processor dependent optimizations
  - Register allocation
  - Peephole optimizations



## Machine Dependent

```
graph TD; MD[Machine Dependent] --> HL[Hide latency]; MD --> MR[Manage resources]; MD --> SF[Special features]; HL --> HL_L["Scheduling<br/>Blocking references<br/>Prefetching<br/>Code layout<br/>Data packing"]; MR --> MR_L["Allocate (registers, tlb slots)<br/>Schedule<br/>Data packing<br/>Coloring memory locations"]; SF --> SF_L["Instruction selection<br/>Peephole optimization"];
```

### Hide latency

Scheduling  
Blocking references  
Prefetching  
Code layout  
Data packing

### Manage resources

Allocate (registers, tlb slots)  
Schedule  
Data packing  
Coloring memory locations

### Special features

Instruction selection  
Peephole optimization

# GCC -O3 [-mtune=generic -march=x86-64] (ver. 4.7.1)

-falign-labels -fasynchronous-unwind-tables -fauto-inc-dec -fbranch-count-reg -fcaller-saves -fcombine-stack-adjustments -fcommon -fcompare-elim -fcprop-registers -fcrossjumping -fcse-follow-jumps -fdebug-types-section -fdefer-pop -fdelete-null-pointer-checks -fdevirtualize -fdwarf2-cfi-asm -fearly-inlining -feliminate-unused-debug-types -fexpensive-optimizations -fforward-propagate -ffunction-cse -fgcse -fgcse-after-reload -fgcse-lm -fgnu-runtime -fguess-branch-probability -fident -fif-conversion -fif-conversion2 -findirect-inlining -finline -finline-atomics -finline-functions -finline-functions-called-once -finline-small-functions -fipa-cp -fipa-cp-clone -fipa-profile -fipa-pure-const -fipa-reference -fipa-sra -fira-share-save-slots -fira-share-spill-slots -fivopts -fkeep-static-consts -fleading-underscore -fmath-errno -fmerge-constants -fmerge-debug-strings -fmove-loop-invariants -fomit-frame-pointer -foptimize-register-move -foptimize-sibling-calls -foptimize-strlen -fpartial-inlining -fpeephole -fpeephole2 -fpredictive-commoning -fprefetch-loop-arrays -free -freg-struct-return -fregmove -freorder-blocks -freorder-functions -frerun-cse-after-loop -fsched-critical-path-heuristic -fsched-dep-count-heuristic -fsched-group-heuristic -fsched-interblock -fsched-last-insn-heuristic -fsched-rank-heuristic -fsched-spec -fsched-spec-insn-heuristic -fsched-stalled-insns-dep -fschedule-insns2 -fshow-column -fshrink-wrap -fsigned-zeros -fsplit-ivs-in-unroller -fsplit-wide-types -fstrict-aliasing -fstrict-overflow -fstrict-volatile-bitfields -fthread-jumps -ftoplevel-reorder -ftrapping-math -ftree-bit-ccp -ftree-builtin-call-dce -ftree-ccp -ftree-ch -ftree-copy-prop -ftree-copyrename -ftree-cselim -ftree-dce -ftree-dominator-opts -ftree-dse -ftree-forwprop -ftree-fre -ftree-loop-distribute-patterns -ftree-loop-if-convert -ftree-loop-im -ftree-loop-ivcanon -ftree-loop-optimize -ftree-parallelize-loops= -ftree-pprop -ftree-pre -ftree-pta -ftree-reassoc -ftree-scev-cprop -ftree-sink -ftree-slp-vectorize -ftree-sra -ftree-switch-conversion -ftree-tail-merge -ftree-ter -ftree-vect-loop-version -ftree-vectorize -ftree-vrp -funit-at-a-time -funswitch-loops -funwind-tables -fvar-tracking -fvar-tracking-assignments -fvect-cost-model -fzero-initialized-in-bss -m128bit-long-double -m64 -m80387 -maccumulate-outgoing-args -malign-stringops -mfancy-math-387 -mfp-ret-in-387 -mglibc -mieee-fp -mmmx -mno-sse4 -mpush-args -mred-zone -msse -msse2 -mtls-direct-seg-refs



# Optimization rule #1

*The First Rule of Program Optimization:*

***Don't do it.***

*The Second Rule (for experts only!):*

***Don't do it yet.***

# False optimizations

## Don't second-guess the compiler

There are things that the compiler can easily optimize by itself, for example, writing `a = a >> 1` should never replace `a = a / 2` if your intention is to divide the variable “a” by 2 and not shifting it. By doing so you are reducing the readability of your code without really improving anything, modern compilers are perfectly able to do that optimization by themselves.

# Outsmarting the Compiler

```
/* original code */
unsigned foo(unsigned char i)
{ // 3*shl, 3*add
    return i + (i<<8) + (i<<16) + (i<<24);
}
```

```
/* attempt to improve foo */
unsigned bar(unsigned char i)
{ // 2*shl, 2*add
    unsigned int j = i + (i << 8);
    return j + (j<<16);
}
```

```
/* "let the compiler do it" */
unsigned baz(unsigned char i)
{ // 1*imul
    return i*0x01010101;
}
```

# Compiler Optimization Obstacles

- Pointer aliasing
  - Pointer arithmetic
  - Global variables
  - Dynamic memory allocation
- Control dependencies
  - Indirect addressing
  - Floating point
- Function calls
  - Side effects
  - External functions
  - Virtual functions
- Optimization barriers
  - `volatile` variables
  - `printf`
  - Intrinsic functions
  - inline assembly
- Correctness overhead
  - Sign-extend
  - Stack overusage
  - Exception handling
  - Hardware bugs

# C/C++ Storage Classes

A storage class defines the scope (visibility) and life-time of variables and/or functions within a C++ Program.

- `auto`
- `register`
- `static`
- `extern`
- `mutable`

# Storage Classes for Optimizers

- `register` - the fastest, the most limited
- `stack` - compiler knows all about them
- `memory` - subject of aliasing
- `extern` - compiler knows nothing about them
- `volatile` - strict barrier

# Pointer Aliasing

- C/C++ language has pointers, C/C++ compiler has its worst problem - **pointer aliasing**
- Pointers **alias** when they point to the same address
  - Writing via one pointer will change the value read through another
- The compiler often doesn't know which pointers alias
  - The compiler must assume that any write through a pointer may affect the value read from any another pointer!
  - This can significantly reduce code efficiency

```
/* what if: timers(t1, t1, t1) */  
void timers(int *t1, int *t2, int *step)  
{  
    *t1 += *step;  
    *t2 += *step;  
}
```



EVERY TIME YOU OPTIMIZE  
POINTER ARITHMETIC

GOD KILLS A KITTEN



# Taking a Variable's Address

```
int N;  
  
getlimit(&N);  
for (i = 0; i < N; i++)  
    ...
```

- Taking the address of a variable means that it must live in memory
  - The variable is then subject to **pointer aliasing**
  - Heavy use of the variable will be costly
- Even if, as in example above, you only take the address once then use it later it's still a **memory-bound variable**
- **Solution**: make a second, non-memory-bound, copy of the variable for intensive use
- **Global variables** are always memory-bound and always subject of aliasing

# More about Variables

## Single responsibility principle

- Every violation of this principle produce false dependency
- Declare variables closer to usage
- You can't improve stack usage by early declaration of variable

# More about Variables

## Variable Storage Classes

- **Register** storage
  - In the best case compiler will be able to assign variable to some register
- **Stack** storage
  - Planned by compiler at compile time
    - No memory allocation happens
    - If you have more variables than registers then part of them are SPILLED to stack
  - Thread-local storage is a kind of stack storage
- **Extern** storage
  - For global and static variables

# More about Variables

Certain data types are more efficient to use for local variables than others. So

Use machine's NATURAL WORD TYPE

- Shorter types require sign-extend or truncation
- Larger types require greater alignment and occupy more registers
- There is only a minimal difference between the efficiency of 32-bit integers and 64-bit integers, as long as you are not doing divisions
- So use 32-bit integers for intermediate values even if input/output has smaller precision

# More about Variables

## Signed vs Unsigned integers

In most cases, there is no difference in speed between using signed and unsigned integers.

But there are a few cases where it matters:

- **Division** by a constant: Unsigned is faster than signed when you divide an integer with a constant. This also applies to the modulo operator %
- **Overflow** behaves differently on signed and unsigned variables
- **Conversion** to floating point might be faster with signed than with unsigned integers

# More about Variables

## Pointers vs References

### REFERENCES ARE POINTERS

with syntactic sugar from the point of compiler view

# More about Variables

## Floating Point

- The main rule: don't mix single and double precision (explicitly add f suffix to constants e.g. 1.0f)
- Single precision float has 2 times smaller memory footprint
- Single precision float better compatible with various coprocessors
- Fixed-point is rarely efficient as replacement of true floating-point calculations

# C/C++ Keywords

- **inline** - language overhead
- **const** - means nothing for optimizer (by default)
- **register** - complete placebo
- **restrict** - ~~C only~~; mostly useless
- **volatile** - optimization barrier



# #define for numeric constants

- `#define CONSTANT 23`
- `const int Constant=23;`
- `static const int Constant=23;`
- `enum { constant=23 };`

No memory references and additions in generated code:

```
void foo(void)
{
    a(constant+3);
    a(CONSTANT+4);
    a(Constant+5);
}
```

```
foo:
```

```
subq $8, %rsp
movl $26, %edi
call a
movl $27, %edi
call a
movl $28, %edi
addq $8, %rsp
jmp a
```

# macro vs inline

```
#define abs(x) ((x)>0?(x):- (x))  
static long abs2(int x) { return x>=0?x:-x; }  
  
int foo(int a) { return abs(a); }  
int bar(int a) { return abs2(a); }
```

Compiler emits inlined branchless code:

```
foo:  
bar:  
    mov     edx,edi  
    sar     edx,31      @ int tmp = x >> (sizeof(x) * 8 - 1);  
    mov     eax,edx  
    xor     eax,edi     @ return (tmp ^ x) - tmp;  
    sub     eax,edx  
    ret
```

# “static” keyword

- Mark constants **static**
  - compiler will not reserve memory slots for static constants
- Mark helper functions **static**
  - **static** gives you internal linkage: compiler will know that it sees all usages of that function
  - compiler will inline static functions used only once
- In C++11 anonymous namespace has exactly the same effect as **static** keyword (it does not in C++98)
- If you have to use global variables then make them **static**

THE END