

Optimization Techniques, Loops

Andrey Kamaev
Itseez UNN 2016

Loop Optimizations

43J0⁴8ε1

Single loop optimizations

Nested loop optimizations

Single Loop Optimizations

- Induction variable optimization
- Loop invariant code motion
- Scalarization
- Loop unswitching
- Loop peeling
- Loop fusion
- Loop fission
- Loop reversal
- Software pipelining
- Loop unrolling
- Loop vectorization

Induction Variable Optimization

- Simplify expressions that change as a linear function of the loop index
 - the loop index is multiplied with a constant
 - replaces a multiplication with an addition
- Often used in address calculations when iterating over an array

```
for (int i = 0; i < N; i++) {  
    int k = 4 * i + m;  
    ...  
}
```

```
int k = m;  
for (int i = 0; i < N; i++) {  
    ...  
    k += 4;  
}
```

Loop Invariant Code Motion

- Eliminate invariant code from loops
 - in particular, try to move expensive calls (`malloc`, `strlen`, file open/close, ...) out of loops

```
for (int i = 0; i < strlen(s); i++)  
{  
    // do something with s[i];  
}
```

```
int len = strlen(s);  
for (int i = 0; i < len; i++)  
{  
    // do something with s[i];  
}
```

Scalarization

- Resolves aliasing conflicts
- Eliminates repetitive memory reads

```
float coeffs[2];

for (int i = 0; i < N; i++) {
    a[i] = b[i]*coeffs[0]+coeffs[1];
}
```

```
float coeffs[2];
float c0 = coeffs[0], c1 = coeffs[1];

for (int i = 0; i < N; i++) {
    a[i] = b[i] * c0 + c1;
}
```

Loop Unswitching

- Move loop invariant conditional constructs out of the loop
 - if or switch statements which are independent of the loop index can be moved outside of the loop
 - the loop is instead repeated in the different branches of the if or case statement
 - removes branch instructions from within the loop, increases ILP

```
for (int i = 0; i < N; i++) {  
    if (a > 0)  
        X[i] = a;  
    else  
        X[i] = 0;  
}
```

```
if (a > 0) {  
    for (int i = 0; i < N; i++)  
        X[i] = a;  
}  
else {  
    for (int i = 0; i < N; i++)  
        X[i] = 0;  
}
```

Loop Peeling

- A small number of iterations from the beginning and/or end of a loop are removed and executed separately
 - for example the handling of boundary conditions
- Removes branches from the loop
 - results in larger basic blocks
 - improves ILP

```
for (int i = 0; i < N; i++) {  
    if (i == 0)  
        X[i] = 0;  
    else if (i == N - 1)  
        X[i] = N;  
    else  
        X[i] = X[i] * c;  
}
```

```
for (int i = 1; i < N - 1; i++) {  
    X[i] = X[i] * c;  
}  
X[N-1] = N;  
X[0] = 0;
```


Loop Fusion

- Loop overhead reduced
- Better instruction overlap
- Temporary locality is improved = lower cache misses
- Be aware of associativity issues with array's mapping to the same cache line

```
for (int i = 0; i < N; i++)  
    x = x * a[i] + b[i];  
for (int j = 0; j < 2*N; j++)  
    y = y + a[j] / b[j];
```

```
for (int i = 0; i < N; i++) {  
    x = x * a[i] + b[i];  
    y = y + a[i] / b[i];  
}  
for (int j = N; j < 2*N; j++)  
    y = y + a[j] / b[j];
```

Loop Fission

- Register pressure reduced
- Cache associativity conflicts are minimized
- Can be used to break loop-carried dependencies

```
for (int i = 0; i < N; i++) {  
    a[i] = b[i] + 1;  
    c[i] = 3 * a[i];  
    f[i] = g[i] + h[i];  
}
```

```
for (int i = 0; i < N; i++) {  
    a[i] = b[i] + 1;  
    c[i] = 3 * a[i];  
}  
for (int j = 0; j < N; j++)  
    f[j] = g[j] + h[j];
```

Loop Reversal

- Change the direction of loop iteration
 - Can improve cache performance
 - Enables other transformations

```
for (int i = 0; i < N; i++) {  
    a[i] = b[i] + 1;  
    c[i] = 3 * a[i];  
}  
for (int j = 0; j < N; j++)  
    d[j] = c[j+1] + 1;
```

```
for (int i = N - 1; i >= 0; i--) {  
    a[i] = b[i] + 1;  
    c[i] = 3 * a[i];  
}  
for (int j = N - 1; j >= 0; j--)  
    d[j] = c[j+1] + 1;
```

Software Pipelining

- Applies instruction scheduling, allowing instructions within a loop to “wrap around” and execute in a different iteration of the loop
- Reduces the impact of long-latency operations, resulting in faster loop execution
- Emulates prefetching of data to reduce the impact of cache misses
- Often used with together with loop unrolling

```
for (int i = 0; i < N; i++) {  
    a[i]++;  
    b[i] = a[i] / 2;  
}
```

```
for (int i = 0; i < N; i++) {  
    a[i] = 3*b[i] + 1;  
}
```

```
if (N > 0) {  
    a[0]++;  
    for (int i = 0; i < N-1; i++) {  
        a[i+1]++;  
        b[i] = a[i] / 2;  
    }  
    b[N-1] = a[N-1] / 2;  
}
```

```
if (N > 0) {  
    int t = b[0];  
    for (int i = 0; i < N-1; i++) {  
        int p = b[i+1];  
        a[i] = 3 * t + 1;  
        t = p;  
    }  
    a[N-1] = 3 * t + 1;  
}
```

Loop Unrolling

- Improves processor pipeline utilization (higher ILP)
- Reduces loop overhead
- Can break loop-carried dependencies
- Needed for vectorization
- Increases register pressure
- Leads to cache associativity conflicts
- Requires additional code to process leftovers
- Can affect results of floating-point arithmetic

```
float sum = 0;
for (int i = 0; i < N; i++) {
    sum += a[i];
}
```

```
float sum = 0;
int i = 0;
for (; i < (N/4)*4; i++) {
    sum += a[i] + a[i+1] + a[i+2] + a[i+3];
}
for (; i < N; i++) {
    sum += a[i];
}
```

```
float sum = 0;
for (int i = 0; i < N; i++) {
    sum += a[i];
}
```

```
float sum1 = 0, sum2 = 0, sum3 = 0, sum4 = 0;
int i = 0;
for (; i < (N/4)*4; i++) {
    sum1 += a[i];
    sum2 += a[i+1];
    sum3 += a[i+2];
    sum4 += a[i+3];
}
float sum = sum1 + sum2 + (sum3 + sum4);
for (; i < N; i++) {
    sum += a[i];
}
```


Nested Loop Optimizations

- Strip Mining
- Loop collapse
- Loop interchange
- Outer loop Unroll and Jam
- Loop blocking (tiling)

Strip Mining

- Usually combined with other optimizations
 - Software prefetch is almost useless without this optimization
- Used to improve locality

```
for (int i = 0; i < N; i++) {  
    a[i] = b[i];  
    c[i] = c[i-1] + a[i];  
}
```

```
for (int i = 0; i < N; i += strip_size) {  
    for (int j = i; j < i + strip_size; j++) {  
        a[j] = b[j];  
        c[j] = c[j-1] + a[j];  
    }  
}
```

Loop Collapse

- Reduces loop overhead
- Reduces number of used registers
- Simplifies vectorization

```
char a[N][N];  
for (int j = 0; j < N; j++)  
    for (int i = 0; i < N; i++)  
        a[i][j]++;
```


```
char a[N * N];  
for (int j = 0; j < N * N; j++)  
    a[j]++;
```

Loop Interchange

- Used mostly to improve spatial locality

```
char a[N][N];  
for (int j = 0; j < N; j++)  
    for (int i = 0; i < N; i++)  
        a[i][j]++;
```

```
char a[N][N];  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        a[i][j]++;
```



Red squares are misses

Outer Loop Unroll and Jam

- More results can be kept in registers or in cache
- Used to reduce number of loads and stores on inner loops with invariants
- Be careful loop body does not become too large
 - Need registers for all intermediate values
 - Increases cache associativity requirements

```
float a[N][N], b[N][N], c[N];  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        a[i][j] = b[i][j] * c[j];
```

```
float a[N][N], b[N][N], c[N];  
for (int i = 0; i < (N/4)*4; i+=4)  
    for (int j = 0; j < N; j++) {  
        a[i+0][j] = b[i+0][j] * c[j];  
        a[i+1][j] = b[i+1][j] * c[j];  
        a[i+2][j] = b[i+2][j] * c[j];  
        a[i+3][j] = b[i+3][j] * c[j];  
    }
```

Loop Blocking


Loop Blocking = Loop Tiling \approx Generalized Strip Mining

- Reduce memory pressure by making use of the caches
- Helps when potential for **re-use is high**
- Hard to choose optimal block size:
 - The bigger the blocks are, the greater the reduction in memory traffic
 - All accessed blocks must be able to fit into cache at the same time
 - If the blocks are too big, the TLB cache will thrash
- Square blocks are simpler, but rectangular blocks, in which the longest dimension corresponds to the inner loop, generally perform the best
- Can help for vectorization
- Can be used to improve precision of floating-point math

```
float A[N][N], B[N][N];
for (int i = 0; i < N; i++) {
    for (int j = 0; j < N; j++) {
        A[i][j] = A[i][j] + B[j][i];
    }
}
```

```
float A[N][N], B[N][N];
for (int i = 0; i < N; i += block_size) {
    for (int j = 0; j < N; j += block_size) {

        for (int ii = i; ii < i + block_size; ii++) {
            for (int jj = j; jj < j + block_size; jj++) {
                A[ii][jj] = A[ii][jj] + B[jj][ii];
            }
        }
    }
}
```



Advanced Examples

Gather-Scatter Optimization

```
for (int i = 1; i < N; i++) {  
    if (t[i] > 0) {  
        a[i] = 2 * b[i-1];  
    }  
}
```

```
int n = 0;  
for (int i = 1; i < N; i++) {  
    if (t[i] > 0) {  
        tmp[n] = i;  
        n++;  
    }  
}  
  
for (int j = 0; j < n; j++) {  
    a[tmp[j]] = 2 * b[tmp[j] - 1];  
}
```

The computationally intensive loop runs only over the indices for which the condition was true and can be better optimized

Ring buffer replacement

This is an example of L1-cache targeted optimization

step #1: scalarization

```
int buf[3];
for (int i = 0; i < N; i++) {
    buf[i%3] = X[i];
    int a = buf[(i-2)%3] * 68;
    a     += buf[(i-1)%3] * 99;
    a     += buf[i%3] * 68;
    X[i] = a;
}
```

```
int b0, b1, b2;
for (int i = 0; i < N; i++) {
    b2 = b1;
    b1 = b0;
    b0 = X[i];
    int a = b2 * 68;
    a     += b1 * 99;
    a     += b0 * 68;
    X[i] = a;
}
```

step #2: unroll 3x

```
int b0, b1, b2;
for (int i = 0; i < N; i++) {
    b2 = b1;
    b1 = b0;
    b0 = X[i];
    int a = b2 * 68;
    a    += b1 * 99;
    a    += b0 * 68;
    X[i] = a;
}
```

```
int b0, b1, b2;
for (int i = 0; i < N; i+=3) {
    b1 = X[3*i];
    int a0 = b0 * 68;
    a0    += b2 * 99;
    a0    += b1 * 68;
    X[3*i] = a0;

    b2 = X[3*i+1];
    int a1 = b1 * 68;
    a1    += b0 * 99;
    a1    += b2 * 68;
    X[3*i+1] = a1;

    b0 = X[3*i+2];
    int a2 = b2 * 68;
    a2    += b1 * 99;
    a2    += b0 * 68;
    X[3*i+2] = a2;
}
```

THE END

Floptimization

Throw in some extra Flops to make the CPU perform more “work” at hardly any extra cost; often there is at least some headroom in the floating-point pipelines when running real applications.

Accumulating something in a register is a classic:

```
for (s = 0.0, i = 0; i < N; ++i) {  
    p[i] = f(q[i]); // actual work  
    s += p[i];      // floptimization  
}
```